

# CSCI461: Introduction to Big Data (Big Data Analytics)

Lab #8: Apache Hadoop & MapReduce



# Agenda

- What is Hadoop? [Lecture Slides]
- Setup the Environment
- Apache Hadoop on Docker using Docker Compose
- HDFS Commands
- Hadoop UI
- MapReduce Example
- Today's Task


# Setup the Environment

- You need to have the following java version:
  - openjdk version "1.8.0\_232" or any similar version [Check [this](#)]
- Install IntelliJ IDEA Community Edition. [Check [this](#)]




# Apache Hadoop on Docker using Docker Compose

- Docker Compose is a tool for defining and running multi-container Docker applications.
- It allows you to define an entire application stack, including services, networks, and volumes, in a single file called docker-compose.yml.
- This file provides a concise and human-readable way to specify how different Docker containers should interact with each other.





```
• • • docker-compose.yml

version: "3.7"
services:
  db:
    image: mysql:8.0.19
    restart: always
    environment:
      - MYSQL_DATABASE=example
      - MYSQL_ROOT_PASSWORD=password
  app:
    build: app
    restart: always
  web:
    build: web
    restart: always
    ports:
      - 80:80
```



build



build



run


run

run




# Apache Hadoop on Docker using Docker Compose

- Download the folder from the following [link](#).
- Go to the directory where the docker-compose.yml exists:
  - Run the following command: `docker-compose up`
- **NOTE:** Make sure you don't have any other related hadoop containers.
- **NOTE:** Make sure you have installed the docker-compose utility if required.



# Apache Hadoop on Docker using Docker Compose

Service	NameNode	DataNode	ResourceManager	NodeManager	HistoryServer
Port	9870	9864	8088	8042	8188



# HDFS Commands

- To verify that hadoop and all daemons are running successfully use:
  - `jps`
- To list all files and directories in HDFS use:
  - `hadoop fs -ls /` or `hdfs dfs -ls /`
- To create a new directory in HDFS use:
  - `hadoop fs -mkdir <dir/>` or `hdfs dfs -mkdir <dir/>`
- To copy data from local to HDFS use:
  - `hadoop fs -put <local-file-path> <hdfs-file-path>` or `hdfs dfs -put <local-file-path> <hdfs-file-path>`
- To copy data from HDFS to local use:
  - `hadoop fs -get <local-file-path> <hdfs-file-path>` or `hdfs dfs -get <local-file-path> <hdfs-file-path>`



# HDFS Commands (CONT.)

- To view data of a file in HDFS use:
  - `hadoop fs -cat <file-path-in-hdfs> or hdfs dfs -cat <file-path-in-hdfs>`
- To move a file from one location to another in HDFS use:
  - `hadoop fs -mv <source-path-in-hdfs> <dest-path-in-hdfs> or hdfs dfs -mv <source-path-in-hdfs> <dest-path-in-hdfs>`
- To copy a file from one location to another in HDFS use:
  - `hadoop fs -cp <source-path-in-hdfs> <dest-path-in-hdfs> or hdfs dfs -cp <source-path-in-hdfs> <dest-path-in-hdfs>`
- To copy data from local to HDFS use:
  - `hadoop fs -copyFromLocal <local-file-path> <hdfs-file-path> or hdfs dfs -copyFromLocal <local-file-path> <hdfs-file-path>`
- To copy data from HDFS to local use:
  - `hadoop fs -copyToLocal <local-file-path> <hdfs-file-path> or hdfs dfs -copyToLocal <local-file-path> <hdfs-file-path>`

# HDFS Commands (CONT.)

- To move data from local to HDFS use:
  - `hadoop fs -moveFromLocal <local-file-path> <hdfs-file-path>` or `hdfs dfs -moveFromLocal <local-file-path> <hdfs-file-path>`
- To move data from HDFS to local use:
  - `hadoop fs -moveToLocal <local-file-path> <hdfs-file-path>` or `hdfs dfs -moveToLocal <local-file-path> <hdfs-file-path>`
- To remove a file in HDFS use:
  - `hadoop fs -rm <file-path-in-hdfs>` or `hdfs dfs -rm <file-path-in-hdfs>`
- To create a file in a specific location in HDFS use:
  - `hadoop fs -touchz <file-path-and-name-in-hdfs>` or `hdfs dfs -touchz <file-path-and-name-in-hdfs>`

**NOTE:** Almost all normal terminal commands are used in HDFS.

# Hands-on Example on MapReduce

- We're going to implement a trivial example on Hadoop utilizing MapReduce for Word Count.
- Assuming we have a cluster of 100 nodes, totaling 100 TB disk space, 10 TB RAM, etc.
- We will read a text file, and we want to produce the number of occurrences of each word.
- Attached with lab slides is the Java code used and how to run the example in addition to the data file used.
- **Let's open IntelliJ.**

# Hands-on Example on MapReduce (CONT.)

- MapReduce jobs consists mainly of 3 classes:
  - **The Driver:** the main class of the job, contains the configurations of the job, submitting the job to the cluster, configuring Mapper and Reducer
  - **The Mapper:** The map method to specify the map behavior, each map deals with one split at a time.
  - **The Reducer:** The reduce method to specify the reduce behavior.
- We will explain the code in IntelliJ.
- Add the following jars to the libraries in IntelliJ: [[Jars](#)]
  - hadoop-common-3.3.1.jar
  - hadoop-mapreduce-client-common-3.3.1.jar
  - hadoop-mapreduce-client-core-3.3.1.jar
  - hadoop-client-api-3.3.1.jar
  - hadoop-client-runtime-3.3.1.jar
  - hadoop-hdfs-3.3.1.jar
- To know more about how to write a MapReduce Job.

# Hands-on Example on MapReduce (CONT.)

- After writing your MapReduce job, we need to have the jar file of the job. [build]
- Then move it to the container. [Same for the data file]
  - `docker cp <file> container-name:<path-in-the-container>`
- Move the data file from the container to the HDFS:
  - `hdfs dfs -copyFromLocal <path-of-data-file> <hdfs-path>`
- Finally run the following commands to submit the job:
  - `hadoop jar <jar-path> <DriverClassNameWithout.java> <input-path-on-hdfs> <output-path-on-hdfs>`

**NOTE:** Output directory shouldn't be created before running the job.

# Today's Task

- You're required to implement the **Inverted Index** algorithm using **MapReduce**.
- Refer to the following [link](#) to understand what is inverted index is?
- On Moodle you're required to submit **ONE** zip file contains:
  - Driver Class (.java)
  - Mapper Class (.java)
  - Reducer Class (.java)
  - Screenshot from your ResourceManager UI. (showing the job)
  - Screenshot from HDFS in terminal/bash. (showing the input and output directories)
  - The output file of running your job.
- **CHEATING = ZERO.**
- DEADLINE: ***Saturday, April 20th 2024 at 11:59 PM.*** (No extension will be made)

# Thanks